

Terrain-Conditioned Cross-Attention for Cross-Region Avalanche Debris Mapping with Sentinel-1 SAR

1

¹Author affiliations withheld for double-blind peer review.

Abstract

Cloud cover, winter darkness, and steep terrain complicate regional mapping of avalanche debris. Sentinel-1 synthetic aperture radar (SAR) provides all-weather observations, but slope-dependent foreshortening, layover, and shadow can resemble bitemporal change. We present TopoRadar-Net, a lightweight change-segmentation model that conditions SAR features on local incidence angle, slope, aspect, elevation, and a fixed-reference aspect coordinate through multiscale cross-attention. We evaluate five architectures on AvalCD, comprising seven events, four mountain systems, and 993 avalanche polygons. Models are trained on Alps and Greenland scenes and tested without site-specific tuning in Tromsø and the Pamirs. Across three seeds, TopoRadar-Net obtains **63.79% $\pm$ 1.88%** pooled macro F1 (mean $\pm$ population standard deviation). Its margin over a bottleneck self-attention U-Net is **+0.38** percentage points ($p = \mathbf{0.869}$), whereas its margin over SiamUNet-diff is **+8.92** percentage points ($p = \mathbf{0.021}$). Ablations show regional rather than uniform benefits: removing aspect conditioning or cross-attention increases pooled F1, while the slope prior overlaps **18.47%** of training positives. A validation-calibrated road-corridor proxy reaches **83.0%** precision and **62.5%** recall in the Pamirs but misses the single Tromsø corridor avalanche in every seed. Terrain conditioning is therefore competitive and diagnostically useful, but not uniformly superior or deployment-ready.

Keywords: Sentinel-1 SAR, snow avalanche mapping, semantic segmentation, radar geometry, cross-region evaluation, deep learning

1 Introduction

Regional avalanche-debris maps support event documentation, hazard assessment, and prioritization of field inspection. Optical observations are often unavailable during major winter storms, while Sentinel-1 C-band synthetic aperture radar (SAR) supplies day-and-night measurements through cloud cover (Eckerstorfer, Malnes, & Müller, 2017; Eckerstorfer, Müller, Malnes, & Oterhals, 2021; Eckerstorfer, Vickers, Malnes, & Grahn, 2019). Avalanche deposits roughen the snow surface and can increase post-event co- and cross-polarized backscatter, making bitemporal change detection useful (Tompkin & Leinss, 2021). AvalCD supplies VV/VH for Livigno, Tromsø, and Pish, but HH/HV for Nuuk. Side-looking geometry also produces slope-dependent foreshortening, layover, shadow, and radiometric variation that can resemble debris.

Recent remote-sensing studies show that deep learning can map slope hazards across diverse environments, but shadows, scale changes, class imbalance, and geographic transfer remain difficult (Chandra, Vaidya, Chauhan, & Shaw, 2025; Nguyen et al., 2025). SAR-specific segmentation and change-image construction introduce additional representation choices (Arjasakusuma & Kusuma, 2025; Wang et al., 2025). In the same hazard domain, Ahmad, Mallick, Hang, Rihan, and Bakshi (2026) model pre-event avalanche susceptibility with attention, whereas our target is post-event debris. Debris mapping for emergency response also needs explicit comparison and operationally meaningful validation rather than accuracy alone (Bayat, Dindar, & Erkal, 2026).

The closest avalanche-segmentation precedents already contain terrain conditioning. Bianchi, Grahn, Eckerstorfer, Malnes, and Vickers (2021) apply a potential-angle-of-reach-conditioned attention mask and slope feature map to Sentinel-1 avalanche segmentation. On AvalCD, Gatti, Mariani, Gallo, and Monti (2026) encode local incidence angle and slope in a separate transformer branch and report that SAR-only

change detection is more consistent than terrain-augmented variants. Our architectural delta is narrower: SAR-change queries attend to pooled terrain keys and values, including a fixed 78° aspect-reference channel. The empirical contribution is the three-seed, leave-region-out characterization of component instability, positive-label conflict in the slope prior, and a bounded corridor proxy—not a claim that terrain conditioning uniformly improves accuracy.

We evaluate whether radar-topographic conditioning improves avalanche change segmentation across unseen mountain systems. TopoRadar-Net uses bitemporal SAR differences as queries and multiscale terrain geometry as keys and values. The study contributes: (1) a compact terrain-conditioned cross-attention architecture; (2) three-seed zero-shot evaluation against four baselines on two held-out mountain systems; (3) spatial-block inference, component ablations, and terrain-stratum stability checks; and (4) construct and corridor audits that expose where the model should not be interpreted operationally. The objective is not to claim uniform superiority, but to establish where terrain conditioning changes performance and where its assumptions fail.

2 Data and Study Design

AvalCD contains co-registered pre/post-event Sentinel-1 intensity, topography, validation masks, and expert-mapped avalanche polygons (Gatti, 2026; Gatti et al., 2026). The seven events span the European Alps, Greenland Arctic, Scandinavian Arctic, and Pamir Mountains (Table 1). Livigno events from April 2024 and January 2025 plus Nuuk 2016 form the training cohort. Livigno March 2025 and Nuuk 2021 are used for validation, checkpoint selection, and decision-threshold calibration. Tromsø 2024 and Pish 2023 remain geographically held out for final testing. This split changes climate, relief, polarization regime, acquisition geometry, and avalanche morphology

Table 1 AvalCD scenes used in the cross-region study. Debris is the number of labelled positive pixels

System	Scene date	Role	Region	Pol.	Polygons	Debris
Alps	3 Apr 2024	Train	Livigno	VV/VH	231	32,606
Alps	29 Jan 2025	Train	Livigno	VV/VH	207	11,438
Alps	18 Mar 2025	Validate	Livigno	VV/VH	79	7,392
Greenland	13 Apr 2016	Train	Nuuk	HH/HV	117	64,822
Greenland	11 Apr 2021	Validate	Nuuk	HH/HV	129	81,595
Pamirs	21 Feb 2023	Test	Pish	VV/VH	113	22,579
Scand. Arctic	20 Dec 2024	Test	Tromsø	VV/VH	117	33,972

simultaneously; results quantify cross-system transfer but do not identify any one changed factor as causal.

The executed loader uses the supplied decibel rasters without percentile scaling. Values outside $[-45, 15]$ dB are marked invalid; after the valid mask is fixed, missing cross- and co-polarized values are replaced by -20 and -12 dB, respectively. The two model positions represent cross/co polarization (VH/VV or HV/HH according to scene), and difference channels retain those positions. The terrain tensor contains local incidence angle (LIA), slope, sine/cosine aspect, elevation from a digital elevation model (DEM), and the fixed coordinate $\cos(\alpha - 78^\circ)$; this last channel is a reference-axis aspect encoding, not authoritative per-event radar heading. Training uses balanced 128×128 patches; evaluation uses complete valid-mask rasters. Physical areas are derived from affine raster transforms rather than full rectangular extents. Figure 1 summarizes the model; Online Resource 1 supplies equations, preprocessing details, complete tables, and additional diagnostics.

3 Method

3.1 Radar-topographic conditioning

Let I^{pre} and I^{post} denote pre- and post-event two-channel cross/co-polarized intensity tensors. A shared residual encoder produces scale- l features F_l^{pre} and F_l^{post} . Their




Fig. 1 TopoRadar-Net schematic. Bitemporal SAR differences query pooled topographic keys and values; a pixelwise terrain gate conditions decoding. The training schematic omits the ground-truth loss input for clarity

difference representation is

$$F_l^\Delta = \text{Conv}[F_l^{\text{post}} - F_l^{\text{pre}}, F_l^{\text{post}}, F_l^{\text{pre}}]. \quad (1)$$

The terrain encoder maps the six-channel geometry tensor to T_l . A Radar-Topographic Cross-Attention Block (RTCAB) forms SAR queries and pooled terrain keys and values,

$$Q = F_l^\Delta W_Q, \quad K = \text{Pool}(T_l) W_K, \quad V = \text{Pool}(T_l) W_V, \quad (2)$$

and computes $A = \text{softmax}(QK^\top / \sqrt{d})V$. A pixelwise gate $\sigma(\text{Conv}(T_l))$ modulates the attended features before progressive decoding. This construction conditions a shift-invariant convolutional representation on spatially varying radar geometry; it does not make the representation physically invariant.

231 3.2 Loss, training, and inference

232
233 The objective combines positive-weighted binary cross-entropy (BCE), Dice loss, and
234
235 a soft slope prior:

$$236 \quad \mathcal{L} = \mathcal{L}_{\text{BCE}} + \mathcal{L}_{\text{Dice}} + 0.5 \mathcal{L}_{\text{slope}}. \quad (3)$$

237
238 The prior discourages probability mass below 5° and above 65° but is treated as
239
240 a heuristic because mapped positives occur in this mask. All models use the same
241
242 optimizer, six-epoch schedule, balanced patches, validation-selected checkpoint, and
243
244 full-scene threshold calibration. We train fixed seeds 42, 123, and 456. Comparators
245 are SiamUNet-diff and SiamUNet-conc (Caye Daudt, Le Saux, & Boulch, 2018), ResU-
246
247 Net (Zhang, Liu, & Wang, 2018), and an internal bottleneck self-attention U-Net. The
248
249 internal baseline uses the same convolutional encoder-decoder scale with four-head
250
251 self-attention only at the deepest feature map; it is neither canonical attention-gated
252
253 U-Net nor Swin-UNet. Parameter counts range from 2.01 to 3.46 million.

254 Inference uses overlapping windows with Hanning blending. Pixel F1 and
255
256 intersection-over-union are computed over valid pixels. Polygon hit rate requires at
257
258 least 30% overlap. A 1,000-draw paired bootstrap resamples 128×128 spatial blocks;
259
260 paired seed-level tests use the three macro-F1 values. The latter have low power and
261
262 are interpreted together with effect sizes and ordering across seeds. No test pixel or
263
264 test label is used for model or threshold selection.

265 For a post-hoc operational proxy, validation-scene component scores calibrate a
266
267 one-feature logistic model and an F2-optimal alert threshold. Frozen test-scene alerts
268
269 are intersected with 30-m buffers around a versioned OpenStreetMap major-road
270
271 snapshot (OpenStreetMap contributors, 2026). This retrospective proxy excludes data
272
273 acquisition, human review, dispatch, and closure decisions. Figures 1 and 2 were
274
275 generated with Python 3.14.7 and Matplotlib 3.11.2.
276

Table 2 Zero-shot test performance across three fixed training seeds (mean $\pm$ descriptive population standard deviation). HR is polygon hit rate at 30% overlap

Method	Tromsø F1	Pish F1	Pooled F1	Tromsø HR
SiamUNet-diff	72.64 $\pm$ 0.39	37.11 $\pm$ 2.44	54.87 $\pm$ 1.03	63.25 $\pm$ 2.52
SiamUNet-conc	78.92 $\pm$ 1.07	47.34 $\pm$ 0.70	63.13 $\pm$ 0.79	82.62 $\pm$ 1.07
ResU-Net	76.60 $\pm$ 1.30	49.25 $\pm$ 3.39	62.92 $\pm$ 1.42	76.92 $\pm$ 4.83
Self-attention U-Net	78.09 $\pm$ 0.99	48.73 $\pm$ 2.74	63.41 $\pm$ 1.03	79.20 $\pm$ 2.45
TopoRadar-Net	77.40 $\pm$ 1.86	50.18 $\pm$ 4.40	63.79 $\pm$ 1.88	75.50 $\pm$ 1.61

3.3 AI-assisted research support

Generative AI tools assisted code refactoring, manuscript language, and checklist-based consistency checking. They did not generate research data or determine final scientific claims. The authors inspected generated material, reran the analysis and verification code, checked citations and numerical statements, and retain full responsibility for the work.

4 Results

4.1 Cross-region performance

TopoRadar-Net records 77.40% $\pm$ 1.86% F1 in Tromsø and 50.18% $\pm$ 4.40% in Pish, giving 63.79% $\pm$ 1.88% pooled macro F1 (Table 2). The bottleneck self-attention U-Net obtains 63.41% $\pm$ 1.03%, SiamUNet-conc 63.13% $\pm$ 0.79%, and ResU-Net 62.92% $\pm$ 1.42%. Their differences from TopoRadar-Net are not significant across three seeds. SiamUNet-diff is lower at 54.87% $\pm$ 1.03%; the paired difference is +8.92 percentage points (95% confidence interval [CI] +3.25 to +14.59, $p = 0.021$). The bottleneck self-attention U-Net leads on seed 123, demonstrating initialization-dependent ordering. Figure 2 visualizes the region-specific ordering.

The 339-block seed-42 bootstrap gives a +3.27-percentage-point micro-F1 difference from the self-attention U-Net (95% CI +1.12 to +5.37, $p = 0.010$), but repeating




Fig. 2 Zero-shot pixel F1 across held-out Tromsø, Pish, and their macro average. Error bars show descriptive population standard deviation over three fixed training seeds

the bootstrap by training seed yields +3.26, −2.99, and +4.16 percentage points. Thus a conditional spatial interval is not a substitute for training uncertainty. On the Nuuk validation scene, all modern architectures lie within 0.63 F1 percentage points, providing additional evidence of parity rather than uniform dominance.

Online Resource 1 reports polygon completeness by European Avalanche Warning Services destructive-size classes ([European Avalanche Warning Services, 2022](#)); these classes are available for Tromsø but not Pish.

4.2 Component sensitivity and failure boundaries

The full architecture is not the best pooled configuration (Table 3). Removing aspect and its fixed-reference coordinate or removing cross-attention increases pooled macro F1 to 65.35% and 65.07%, respectively, mainly through gains in Pish. In Tromsø, seed-42 paired block tests favour the full model over No-LIA by 2.04 percentage points ($p < 0.001$) and No-CrossAttn by 1.42 percentage points ($p = 0.010$). These comparisons establish component sensitivity, not a causal terrain-correction mechanism.

Table 3 Ablation performance across three seeds. The highest pooled value is obtained without aspect and its fixed-reference coordinate

Configuration	Tromsø F1	Pish F1	Pooled F1
Full	77.40 ± 1.86	50.18 ± 4.40	63.79 ± 1.88
No local incidence angle	75.67 ± 1.76	52.13 ± 1.11	63.90 ± 1.37
No cross-attention	76.85 ± 0.80	53.30 ± 0.93	65.07 ± 0.53
No slope prior	76.70 ± 0.91	51.99 ± 1.31	64.35 ± 1.05
No aspect/reference coordinate	77.38 ± 2.66	53.32 ± 0.42	65.35 ± 1.54

The slope-prior mask contains 18.47% of valid training positives, 0.95% of validation positives, and 0.87% of held-out positives. Within the held-out penalty stratum, the full model has lower three-seed positive recall than No-GeoLoss in Tromsø ($50.47\% \pm 3.18\%$ versus $71.29\% \pm 6.39\%$) and Pish ($1.08\% \pm 1.52\%$ versus $6.45\% \pm 6.97\%$). The term is therefore a precision–recall prior, not a physical-admissibility rule.

Seed-42 terrain strata initially suggest positive backslope differences, but these do not persist across seeds: TopoRadar-Net minus the self-attention U-Net averages $-0.34\% \pm 6.28\%$ in Tromsø and $+1.20\% \pm 7.73\%$ in Pish. Extreme-steep terrain is highly variable and both models perform poorly in the small Pish shadow stratum. Complete stratum results are in Online Resource 1.

4.3 Retrospective corridor proxy

The affine valid-mask area is 196.820 km^2 in Tromsø and 267.435 km^2 in Pish. In Tromsø, TopoRadar-Net produces 74.77 ha of false positives at 0.38 ha km^{-2} . In Pish, it produces 129.69 ha at 0.48 ha km^{-2} . Relative to the self-attention U-Net, this is a 30.0% reduction in Pish (129.69 versus 185.32 ha) but a 28.5% increase in Tromsø (74.77 versus 58.19 ha). Validation-calibrated road-corridor alerts achieve 83.0% precision and 62.5% recall in Pish. The single road-intersecting reference avalanche in Tromsø is missed by every TopoRadar-Net seed. All cached complete-scene runs are

below 5 s on Apple Metal Performance Shaders (MPS) hardware, but timing excludes upstream acquisition and human review.

5 Discussion

TopoRadar-Net is competitive across unseen mountain systems, but the strongest baselines are statistically indistinguishable across three seeds. This result is important for remote-sensing model design: adding physically interpretable variables does not guarantee a better predictor. Recent hazard-mapping studies similarly motivate multiscale or attention modules across diverse environments (Chandra et al., 2025; Nguyen et al., 2025), while our cross-region ablations show that auxiliary geometry can reverse model ordering. The contribution is therefore the controlled characterization of terrain conditioning under geographic shift, including evidence against a uniform advantage.

The construct audit changes the interpretation of the slope loss. Genuine positives inside the penalty mask arise through runout, raster alignment, and terrain-model uncertainty; suppressing them can increase precision while reducing recall. A hard “impossible terrain” label would be unsupported. Likewise, seed-42 terrain-stratum differences are descriptive because seed repetition does not preserve their magnitude or sign.

Operationally, the component and corridor analyses resemble recent efforts to connect mapped debris with transport access (Bayat et al., 2026), but they remain retrospective. The Pish corridor result is encouraging; the Tromsø miss falsifies any claim of region-independent completeness. Prospective evaluation with warning services or road authorities must measure acquisition delay, review burden, missed-event cost, and decision outcomes before deployment.

Limitations include seven events from one benchmark, three training seeds, correlated spatial blocks, threshold calibration on only two scenes, incomplete event

metadata for size analysis outside Tromsø, and the non-equivariant directional aug-
mentation inherited by the released checkpoints. Sentinel-1 C-band also has an
irreducible shadow boundary in very steep relief. Future validation in the Indian
Himalaya would be valuable, but this study contains no Indian test data and makes
no India-specific accuracy claim. Application to RISAT or NISAR would require
sensor-specific calibration and independent Himalayan labels.

6 Conclusion

Terrain-conditioned cross-attention provides competitive Sentinel-1 avalanche-debris
segmentation and useful failure diagnostics, but it is not uniformly superior to stan-
dard change networks. TopoRadar-Net reaches 63.79% pooled macro F1 across two
unseen mountain systems; only its margin over SiamUNet-diff is significant across
seeds. Aspect and cross-attention ablations improve pooled F1, and the slope prior
suppresses genuine positives. The corridor proxy succeeds in Pish but fails on the
only Tromsø corridor event. The defensible conclusion is therefore bounded: radar-
topographic conditioning can alter regional precision and recall, while cross-region
reliability still requires explicit validation.

Data Availability/Supplementary Information

The AvalCD benchmark is publicly available at <https://doi.org/10.5281/zenodo.15863589>. **Online Resource 1.** Extended methods, complete statistical and sub-
group results, and reproducibility identifiers. Code and checkpoint files are publicly
archived; the identifying repository URL is withheld from the anonymous manuscript
and supplied separately to the editorial office for double-blind review.

Statements and Declarations

Funding. No external grant funding was received for this study.

507 **Competing interests.** The authors declare no competing financial or non-
508 financial interests.
509

510 **Ethics approval and consent.** Not applicable; the study uses public remote-
511 sensing data and contains no human participants or animals.
512

513 **Consent for publication.** Not applicable.
514

515 **Author contributions.** The complete CRediT statement is provided on the
516 separate title page.
517

518 **AI disclosure.** Generative AI assistance is documented in the Methods section.
519 The authors retain responsibility for all code, analyses, citations, and text.
520
521

522 523 **References** 524

525
526 Ahmad, I., Mallick, J., Hang, H.T., Rihan, M., Bakshi, A. (2026). An explainable
527 bayesian-optimized hybrid deep learning framework with attention mechanisms
528 for modeling snow avalanche susceptibility. *Journal of the Indian Society of*
529 *Remote Sensing*, 54(7), 2749–2773, [https://doi.org/10.1007/s12524-026-02459](https://doi.org/10.1007/s12524-026-02459-1)
530
531
532 -1
533

534
535
536
537 Arjasakusuma, S., & Kusuma, S.S. (2025). Assessment of deep learning based image
538 segmentation for identifying floating net cages from very high-resolution capella
539 synthetic aperture radar (SAR) data. *Journal of the Indian Society of Remote*
540 *Sensing*, 53(2), 485–499, <https://doi.org/10.1007/s12524-024-02022-w>
541
542
543

544
545
546 Bayat, Z., Dindar, S., Erkal, B.G. (2026). Deep learning for post-earthquake road
547 blockage and debris detection in resource-limited contexts. *Journal of the Indian*
548 *Society of Remote Sensing*, 54(3), 1257–1285, [https://doi.org/10.1007/s12524](https://doi.org/10.1007/s12524-025-02394-7)
549
550
551 -025-02394-7
552

- Bianchi, F.M., Grahm, J., Eckerstorfer, M., Malnes, E., Vickers, H. (2021). Snow
avalanche segmentation in SAR images with fully convolutional neural networks.
*IEEE Journal of Selected Topics in Applied Earth Observations and Remote
Sensing*, 14, 75–82, <https://doi.org/10.1109/JSTARS.2020.3036914>
- Caye Daudt, R., Le Saux, B., Boulch, A. (2018). Fully convolutional siamese networks
for change detection. *2018 25th IEEE International Conference on Image Processing
(ICIP)* (pp. 4063–4067).
- Chandra, N., Vaidya, H., Chauhan, P., Shaw, R. (2025). *YOLOv5 based on atten-
tion mechanism for landslide information extraction from multi-source remote
sensing images of diverse geological environments*. Advance online publication
in the Journal of the Indian Society of Remote Sensing.
- Eckerstorfer, M., Malnes, E., Müller, K. (2017). A complete snow avalanche activity
record from a norwegian forecasting region using Sentinel-1 satellite-radar data.
Cold Regions Science and Technology, 144, 39–51, [https://doi.org/10.1016/
j.coldregions.2017.08.004](https://doi.org/10.1016/j.coldregions.2017.08.004)
- Eckerstorfer, M., Müller, K., Malnes, E., Oterhals, H.D. (2021). Manual and automatic
detection of dry snow avalanches in Sentinel-1 SAR images. *Ieee international
geoscience and remote sensing symposium (igarss)* (pp. 247–250).
- Eckerstorfer, M., Vickers, H., Malnes, E., Grahm, J. (2019). Near-real time auto-
matic snow avalanche activity monitoring system using Sentinel-1 SAR data in
Norway. *Remote Sensing*, 11(23), 2863, <https://doi.org/10.3390/rs11232863>

European Avalanche Warning Services (2022). *Avalanche size classification standards* (classes D1 through D5). <https://www.avalanches.org/standards/avalanche-size/>. (Accessed: September 2026)

Gatti, M. (2026). *AvalCD*. Zenodo.

Gatti, M., Mariani, A., Gallo, I., Monti, F. (2026). *Large-scale avalanche mapping from SAR images with deep learning-based change detection*. arXiv preprint arXiv:2603.22658.

Nguyen, H., Dang, K.B., Giang, T.L., Dang, V.B., Dang, N.V., Dao, M.D. (2025). *Exploring various deep learning models for high-precision landslide tracing in very-high resolution remote sensing imagery*. Advance online publication in the Journal of the Indian Society of Remote Sensing.

OpenStreetMap contributors (2026). *OpenStreetMap*. <https://www.openstreetmap.org/copyright>. (Major-road snapshot acquired through the Overpass API on 19 September 2026 under ODbL 1.0)

Tompkin, C., & Leinss, S. (2021). Backscatter characteristics of snow avalanches for mapping with local resolution weighting. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 14, 4452–4464, <https://doi.org/10.1109/JSTARS.2021.3074418>

Wang, S., Wang, D., Shi, J., Zhang, Z., Wang, H., Li, X., Guo, Y. (2025). Unsupervised SAR image change detection method based on fused difference images. *Journal of the Indian Society of Remote Sensing*, 53(12), 4373–4382, <https://doi.org/10.1007/s12524-025-02259-z>

Zhang, Z., Liu, Q., Wang, Y. (2018). Road extraction by deep residual U-Net. 645
IEEE Geoscience and Remote Sensing Letters, 15(5), 749–753, [https://doi.org/](https://doi.org/10.1109/LGRS.2018.2802944) 646
[10.1109/LGRS.2018.2802944](https://doi.org/10.1109/LGRS.2018.2802944) 647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690