

Terrain-Conditioned Cross-Attention for Cross-Region Avalanche Debris Mapping with Sentinel-1 SAR

Abstract

Cloud cover, winter darkness, and steep terrain complicate regional mapping of avalanche debris. Sentinel-1 synthetic aperture radar (SAR) provides all-weather observations, but slope-dependent foreshortening, layover, and shadow can resemble bitemporal change. We present TopoRadar-Net, a lightweight change-segmentation model that conditions SAR features on local incidence angle, slope, aspect, elevation, and a fixed-reference aspect coordinate through multiscale cross-attention. We evaluate five architectures on AvalCD, comprising seven events, four mountain systems, and 993 avalanche polygons. Models are trained on Alps and Greenland scenes and tested without site-specific tuning in Tromsø and the Pamirs. Across three seeds, TopoRadar-Net obtains $63.79\% \pm 1.88\%$ pooled macro F1 (mean $\pm$ population standard deviation). Its margin over a bottleneck self-attention U-Net is +0.38 percentage points ($p = 0.869$), whereas its margin over SiamUNet-diff is +8.92 percentage points ($p = 0.021$). Ablations show regional rather than uniform benefits: removing aspect conditioning or cross-attention increases pooled F1, while the slope prior overlaps 18.47% of training positives. A validation-calibrated road-corridor proxy reaches 83.0% precision and 62.5% recall in the Pamirs but misses the single Tromsø corridor avalanche in every seed. Terrain conditioning is therefore competitive and diagnostically useful, but not uniformly superior or deployment-ready.

Keywords

Sentinel-1 SAR, snow avalanche mapping, semantic segmentation, radar geometry, cross-region evaluation, deep learning

25 **Introduction**

26 Regional avalanche-debris maps support event documentation, hazard assessment, and prioritization
27 of field inspection. Optical observations are often unavailable during major winter storms, while
28 Sentinel-1 C-band synthetic aperture radar (SAR) supplies day-and-night measurements through
29 cloud cover (Eckerstorfer et al., 2017; Eckerstorfer et al., 2019; Eckerstorfer et al., 2021).

30 Avalanche deposits roughen the snow surface and can increase post-event co- and cross-polarized
31 backscatter, making bitemporal change detection useful (Tompkin & Leinss, 2021). AvalCD
32 supplies VV/VH for Livigno, Tromsø, and Pish, but HH/HV for Nuuk. Side-looking geometry also
33 produces slope-dependent foreshortening, layover, shadow, and radiometric variation that can
34 resemble debris.

35 Recent remote-sensing studies show that deep learning can map slope hazards across diverse
36 environments, but shadows, scale changes, class imbalance, and geographic transfer remain difficult
37 (Chandra et al., 2025; Nguyen et al., 2025). SAR-specific segmentation and change-image
38 construction introduce additional representation choices (Arjasakusuma & Kusuma, 2025; Wang et
39 al., 2025). In the same hazard domain, Ahmad et al. (2026) model pre-event avalanche
40 susceptibility with attention, whereas our target is post-event debris. Debris mapping for emergency
41 response also needs explicit comparison and operationally meaningful validation rather than
42 accuracy alone (Bayat et al., 2026).

43 The closest avalanche-segmentation precedents already contain terrain conditioning. Bianchi et al.
44 (2021) apply a potential-angle-of-reach-conditioned attention mask and slope feature map to
45 Sentinel-1 avalanche segmentation. On AvalCD, Gatti et al. (2026) encode local incidence angle
46 and slope in a separate transformer branch and report that SAR-only change detection is more
47 consistent than terrain-augmented variants. Our architectural delta is narrower: SAR-change queries
48 attend to pooled terrain keys and values, including a fixed 78° aspect-reference channel. The
49 empirical contribution is the three-seed, leave-region-out characterization of component instability,

50 positive-label conflict in the slope prior, and a bounded corridor proxy—not a claim that terrain
51 conditioning uniformly improves accuracy.

52 We evaluate whether radar-topographic conditioning improves avalanche change segmentation
53 across unseen mountain systems. TopoRadar-Net uses bitemporal SAR differences as queries and
54 multiscale terrain geometry as keys and values. The study contributes: (1) a compact terrain-
55 conditioned cross-attention architecture; (2) three-seed zero-shot evaluation against four baselines
56 on two held-out mountain systems; (3) spatial-block inference, component ablations, and terrain-
57 stratum stability checks; and (4) construct and corridor audits that expose where the model should
58 not be interpreted operationally. The objective is not to claim uniform superiority, but to establish
59 where terrain conditioning changes performance and where its assumptions fail.

60 **Data and Study Design**

61 AvalCD contains co-registered pre/post-event Sentinel-1 intensity, topography, valid-data masks,
62 and expert-mapped avalanche polygons (Gatti et al., 2026; Gatti, 2026). The seven events span the
63 European Alps, Greenland Arctic, Scandinavian Arctic, and Pamir Mountains (Table 1). Livigno
64 events from April 2024 and January 2025 plus Nuuk 2016 form the training cohort. Livigno March
65 2025 and Nuuk 2021 are used for validation, checkpoint selection, and decision-threshold
66 calibration. Tromsø 2024 and Pish 2023 remain geographically held out for final testing. This split
67 changes climate, relief, polarization regime, acquisition geometry, and avalanche morphology
68 simultaneously; results quantify cross-system transfer but do not identify any one changed factor as
69 causal.

70 **Table 1.** *AvalCD scenes used in the cross-region study. Debris is the number of labelled positive*
71 *pixels*

System	Scene date	Role	Region	Pol.	Polygons	Debris
Alps	3 Apr 2024	Train	Livigno	VV/VH	231	32,606
Alps	29 Jan 2025	Train	Livigno	VV/VH	207	11,438
Alps	18 Mar 2025	Validate	Livigno	VV/VH	79	7,392
Greenland	13 Apr 2016	Train	Nuuk	HH/HV	117	64,822
Greenland	11 Apr 2021	Validate	Nuuk	HH/HV	129	81,595
Pamirs	21 Feb 2023	Test	Pish	VV/VH	113	22,579

System	Scene date	Role	Region	Pol.	Polygons	Debris
Scand. Arctic	20 Dec 2024	Test	Tromsø	VV/VH	117	33,972

The executed loader uses the supplied decibel rasters without percentile scaling. Values outside $[-45, 15]$ dB are marked invalid; after the valid mask is fixed, missing cross- and co-polarized values are replaced by -20 and -12 dB, respectively. The two model positions represent cross/co polarization (VH/VV or HV/HH according to scene), and difference channels retain those positions. The terrain tensor contains local incidence geometry, slope, sine/cosine aspect, elevation, and the fixed coordinate $\cos(\alpha - 78^\circ)$; this last channel is a reference-axis aspect encoding, not authoritative per-event radar heading. Training uses balanced 128×128 patches; evaluation uses complete valid-mask rasters. Physical areas are derived from affine raster transforms rather than full rectangular extents. Figure 1 summarizes the model; Online Resource 1 supplies equations, preprocessing details, complete tables, and additional diagnostics.

Method

Radar-topographic conditioning

Let I^{pre} and I^{post} denote pre- and post-event two-channel cross/co-polarized intensity tensors. A shared residual encoder produces scale- l features F_l^{pre} and F_l^{post} . Their difference representation is

$$F_l^\Delta = \text{Conv}[F_l^{post} - F_l^{pre}, F_l^{post}, F_l^{pre}].$$

The terrain encoder maps the six-channel geometry tensor to T_l . A Radar-Topographic Cross-Attention Block forms SAR queries and pooled terrain keys and values,

$$Q = F_l^\Delta W_Q, K = \text{Pool}(T_l) W_K, V = \text{Pool}(T_l) W_V,$$

and computes $A = \text{softmax}(QK^\top / \sqrt{d})V$. A pixelwise gate $\sigma(\text{Conv}(T_l))$ modulates the attended features before progressive decoding. This construction conditions a shift-invariant convolutional representation on spatially varying radar geometry; it does not make the representation physically invariant.




Fig. 1. TopoRadar-Net schematic. Bitemporal SAR differences query pooled topographic keys and values; a pixelwise terrain gate conditions decoding. The training schematic omits the ground-truth loss input for clarity

Loss, training, and inference

The objective combines positive-weighted binary cross-entropy, Dice loss, and a soft slope prior:

$$L = L_{BCE} + L_{Dice} + 0.5 L_{slope}.$$

The prior discourages probability mass below 5° and above 65° but is treated as a heuristic because mapped positives occur in this mask. All models use the same optimizer, six-epoch schedule, balanced patches, validation-selected checkpoint, and full-scene threshold calibration. We train fixed seeds 42, 123, and 456. Comparators are SiamUNet-diff and SiamUNet-conc (Caye Daudt et al., 2018), ResU-Net (Zhang et al., 2018), and an internal bottleneck self-attention U-Net. The internal baseline uses the same convolutional encoder–decoder scale with four-head self-attention only at the deepest feature map; it is neither canonical attention-gated U-Net nor Swin-UNet. Parameter counts range from 2.01 to 3.46 million.

Inference uses overlapping windows with Hanning blending. Pixel F1 and intersection-over-union are computed over valid pixels. Polygon hit rate requires at least 30 % overlap. A 1,000-draw paired

bootstrap resamples 128×128 spatial blocks; paired seed-level tests use the three macro-F1 values. The latter have low power and are interpreted together with effect sizes and ordering across seeds. No test pixel or test label is used for model or threshold selection.

For a post-hoc operational proxy, validation-scene component scores calibrate a one-feature logistic model and an F2-optimal alert threshold. Frozen test-scene alerts are intersected with 30-m buffers around a versioned OpenStreetMap major-road snapshot (OpenStreetMap contributors, 2026). This retrospective proxy excludes data acquisition, human review, dispatch, and closure decisions.

AI-assisted research support

Generative AI tools assisted code refactoring, manuscript language, and checklist-based consistency checking. They did not generate research data or determine final scientific claims. The authors inspected generated material, reran the analysis and verification code, checked citations and numerical statements, and retain full responsibility for the work.

Results


Cross-region performance

TopoRadar-Net records $77.40\% \pm 1.86\%$ F1 in Tromsø and $50.18\% \pm 4.40\%$ in Pish, giving $63.79\% \pm 1.88\%$ pooled macro F1 (Table 2). The bottleneck self-attention U-Net obtains $63.41\% \pm 1.03\%$, SiamUNet-conc $63.13\% \pm 0.79\%$, and ResU-Net $62.92\% \pm 1.42\%$. Their differences from TopoRadar-Net are not significant across three seeds. SiamUNet-diff is lower at $54.87\% \pm 1.03\%$; the paired difference is +8.92 percentage points (95% CI +3.25 to +14.59, $p=0.021$). The bottleneck self-attention U-Net leads on seed 123, demonstrating initialization-dependent ordering. Figure 2 visualizes the region-specific ordering.

Table 2. Zero-shot test performance across three fixed training seeds (mean $\pm$ descriptive population standard deviation). HR is polygon hit rate at 30% overlap

Method	Tromsø F1	Pish F1	Pooled F1	Tromsø HR
SiamUNet-diff	72.64 ± 0.39	37.11 ± 2.44	54.87 ± 1.03	63.25 ± 2.52
SiamUNet-conc	78.92 ± 1.07	47.34 ± 0.70	63.13 ± 0.79	82.62 ± 1.07

Method	Tromsø F1	Pish F1	Pooled F1	Tromsø HR
ResU-Net	76.60 ± 1.30	49.25 ± 3.39	62.92 ± 1.42	76.92 ± 4.83
Self-attention U-Net	78.09 ± 0.99	48.73 ± 2.74	63.41 ± 1.03	79.20 ± 2.45
TopoRadar-Net	77.40 ± 1.86	50.18 ± 4.40	63.79 ± 1.88	75.50 ± 1.61



134

135 **Fig. 2.** Zero-shot pixel F1 across held-out Tromsø, Pish, and their macro average. Error bars show
136 descriptive population standard deviation over three fixed training seeds

137 The 339-block seed-42 bootstrap gives a +3.27-percentage-point micro-F1 difference from the self-
138 attention U-Net (95% CI +1.12 to +5.37, $p=0.010$), but repeating the bootstrap by training seed
139 yields +3.26, -2.99 , and $+4.16$ percentage points. Thus a conditional spatial interval is not a
140 substitute for training uncertainty. On the Nuuk validation scene, all modern architectures lie within
141 0.63 F1 percentage points, providing additional evidence of parity rather than uniform dominance.

142 Online Resource 1 reports polygon completeness by European Avalanche Warning Services
143 destructive-size classes (European Avalanche Warning Services, 2022); these classes are available
144 for Tromsø but not Pish.

145 Component sensitivity and failure boundaries

146 The full architecture is not the best pooled configuration (Table 3). Removing aspect and its fixed-
147 reference coordinate or removing cross-attention increases pooled macro F1 to 65.35% and 65.07%,
148 respectively, mainly through gains in Pish. In Tromsø, seed-42 paired block tests favour the full
149 model over No-LIA by 2.04 percentage points ($p < 0.001$) and No-CrossAttn by 1.42 percentage
150 points ($p = 0.010$). These comparisons establish component sensitivity, not a causal terrain-correction
151 mechanism.

152 **Table 3.** Ablation performance across three seeds. The highest pooled value is obtained without
153 aspect and its fixed-reference coordinate

Configuration	Tromsø F1	Pish F1	Pooled F1
Full	77.40 ± 1.86	50.18 ± 4.40	63.79 ± 1.88
No local incidence angle	75.67 ± 1.76	52.13 ± 1.11	63.90 ± 1.37
No cross-attention	76.85 ± 0.80	53.30 ± 0.93	65.07 ± 0.53
No slope prior	76.70 ± 0.91	51.99 ± 1.31	64.35 ± 1.05
No aspect/reference coordinate	77.38 ± 2.66	53.32 ± 0.42	65.35 ± 1.54

154 The slope-prior mask contains 18.47 % of valid training positives, 0.95 % of validation positives,
155 and 0.87 % of held-out positives. Within the held-out penalty stratum, the full model has lower
156 three-seed positive recall than No-GeoLoss in Tromsø ($50.47 \% \pm 3.18 \%$ versus $71.29 \% \pm 6.39 \%$)
157 and Pish ($1.08 \% \pm 1.52 \%$ versus $6.45 \% \pm 6.97 \%$). The term is therefore a precision–recall prior, not
158 a physical-admissibility rule.

159 Seed-42 terrain strata initially suggest positive backslope differences, but these do not persist across
160 seeds: TopoRadar-Net minus the self-attention U-Net averages $-0.34 \% \pm 6.28 \%$ in Tromsø and
161 $+1.20 \% \pm 7.73 \%$ in Pish. Extreme-steep terrain is highly variable and both models perform poorly
162 in the small Pish shadow stratum. Complete stratum results are in Online Resource 1.

163 Retrospective corridor proxy

164 The affine valid-mask area is 196.820 km² in Tromsø and 267.435 km² in Pish. In Tromsø,
165 TopoRadar-Net produces 74.77 ha of false positives at 0.38 ha km⁻². In Pish, it produces 129.69 ha

166 at 0.48 ha km⁻². Relative to the self-attention U-Net, this is a 30.0% reduction in Pish (129.69
167 versus 185.32 ha) but a 28.5% increase in Tromsø (74.77 versus 58.19 ha). Validation-calibrated
168 road-corridor alerts achieve 83.0% precision and 62.5% recall in Pish. The single road-intersecting
169 reference avalanche in Tromsø is missed by every TopoRadar-Net seed. All cached complete-scene
170 runs are below 5 s on Apple MPS hardware, but timing excludes upstream acquisition and human
171 review.

172 **Discussion**

173 TopoRadar-Net is competitive across unseen mountain systems, but the strongest baselines are
174 statistically indistinguishable across three seeds. This result is important for remote-sensing model
175 design: adding physically interpretable variables does not guarantee a better predictor. Recent
176 hazard-mapping studies similarly motivate multiscale or attention modules across diverse
177 environments (Chandra et al., 2025; Nguyen et al., 2025), while our cross-region ablations show
178 that auxiliary geometry can reverse model ordering. The contribution is therefore the controlled
179 characterization of terrain conditioning under geographic shift, including evidence against a
180 uniform advantage.

181 The construct audit changes the interpretation of the slope loss. Genuine positives inside the penalty
182 mask arise through runout, raster alignment, and terrain-model uncertainty; suppressing them can
183 increase precision while reducing recall. A hard “impossible terrain” label would be unsupported.
184 Likewise, seed-42 terrain-stratum differences are descriptive because seed repetition does not
185 preserve their magnitude or sign.

186 Operationally, the component and corridor analyses resemble recent efforts to connect mapped
187 debris with transport access (Bayat et al., 2026), but they remain retrospective. The Pish corridor
188 result is encouraging; the Tromsø miss falsifies any claim of region-independent completeness.
189 Prospective evaluation with warning services or road authorities must measure acquisition delay,
190 review burden, missed-event cost, and decision outcomes before deployment.

191 Limitations include seven events from one benchmark, three training seeds, correlated spatial
192 blocks, threshold calibration on only two scenes, incomplete event metadata for size analysis
193 outside Tromsø, and the non-equivariant directional augmentation inherited by the released
194 checkpoints. Sentinel-1 C-band also has an irreducible shadow boundary in very steep relief. Future
195 validation in the Indian Himalaya would be valuable, but this study contains no Indian test data and
196 makes no India-specific accuracy claim. Application to RISAT or NISAR would require sensor-
197 specific calibration and independent Himalayan labels.

198 **Conclusion**

199 Terrain-conditioned cross-attention provides competitive Sentinel-1 avalanche-debris segmentation
200 and useful failure diagnostics, but it is not uniformly superior to standard change networks.
201 TopoRadar-Net reaches 63.79% pooled macro F1 across two unseen mountain systems; only its
202 margin over SiamUNet-diff is significant across seeds. Aspect and cross-attention ablations
203 improve pooled F1, and the slope prior suppresses genuine positives. The corridor proxy succeeds
204 in Pish but fails on the only Tromsø corridor event. The defensible conclusion is therefore bounded:
205 radar-topographic conditioning can alter regional precision and recall, while cross-region reliability
206 still requires explicit validation.

207 **Data Availability/Supplementary Information**

208 The AvalCD benchmark is publicly available at <https://doi.org/10.5281/zenodo.15863589>. Online
209 Resource 1 contains extended methods, complete statistical and subgroup results, and
210 reproducibility identifiers. Code and checkpoint files are publicly archived; the identifying
211 repository URL is withheld from the anonymous manuscript and supplied separately to the editorial
212 office for double-blind review.

213 **Statements and Declarations**

214 **Funding.** No external grant funding was received for this study.

215 **Competing interests.** The authors declare no competing financial or non-financial interests.

216 **Ethics approval and consent.** Not applicable; the study uses public remote-sensing data and
217 contains no human participants or animals.

218 **Consent for publication.** Not applicable.

219 **Author contributions.** The complete CRediT statement is provided on the separate title page.

220 **AI disclosure.** Generative AI assistance is documented in the Methods section. The authors retain
221 responsibility for all code, analyses, citations, and text.

222 Ahmad, I., Mallick, J., Hang, H. T., Rihan, M., & Bakshi, A. (2026). An explainable bayesian-
223 optimized hybrid deep learning framework with attention mechanisms for modeling snow
224 avalanche susceptibility. *Journal of the Indian Society of Remote Sensing*, 54(7), 2749–2773.
225 <https://doi.org/10.1007/s12524-026-02459-1>

226 Arjasakusuma, S., & Kusuma, S. S. (2025). Assessment of deep learning based image segmentation
227 for identifying floating net cages from very high-resolution capella synthetic aperture radar (SAR)
228 data. *Journal of the Indian Society of Remote Sensing*, 53(2), 485–499.
229 <https://doi.org/10.1007/s12524-024-02022-w>

230 Bayat, Z., Dindar, S., & Erkal, B. G. (2026). Deep learning for post-earthquake road blockage and
231 debris detection in resource-limited contexts. *Journal of the Indian Society of Remote Sensing*,
232 54(3), 1257–1285. <https://doi.org/10.1007/s12524-025-02394-7>

233 Bianchi, F. M., Grahn, J., Eckerstorfer, M., Malnes, E., & Vickers, H. (2021). Snow avalanche
234 segmentation in SAR images with fully convolutional neural networks. *IEEE Journal of Selected*
235 *Topics in Applied Earth Observations and Remote Sensing*, 14, 75–82.
236 <https://doi.org/10.1109/JSTARS.2020.3036914>

237 Caye Daudt, R., Le Saux, B., & Boulch, A. (2018). Fully convolutional siamese networks for
238 change detection. *2018 25th IEEE International Conference on Image Processing (ICIP)*, 4063–
239 4067. <https://doi.org/10.1109/ICIP.2018.8451652>

240 Chandra, N., Vaidya, H., Chauhan, P., & Shaw, R. (2025). YOLOv5 based on attention mechanism
 241 for landslide information extraction from multi-source remote sensing images of diverse geological
 242 environments. *Journal of the Indian Society of Remote Sensing*, 53, 1–21.
 243 <https://doi.org/10.1007/s12524-025-02251-7>

244 Eckerstorfer, M., Malnes, E., & Müller, K. (2017). A complete snow avalanche activity record from
 245 a norwegian forecasting region using Sentinel-1 satellite-radar data. *Cold Regions Science and*
 246 *Technology*, 144, 39–51. <https://doi.org/10.1016/j.coldregions.2017.08.004>

247 Eckerstorfer, M., Müller, K., Malnes, E., & Oterhals, H. D. (2021). Manual and automatic detection
 248 of dry snow avalanches in Sentinel-1 SAR images. *IEEE International Geoscience and Remote*
 249 *Sensing Symposium (IGARSS)*, 247–250. <https://doi.org/10.1109/IGARSS47720.2021.9554370>

250 Eckerstorfer, M., Vickers, H., Malnes, E., & Grahn, J. (2019). Near-real time automatic snow
 251 avalanche activity monitoring system using Sentinel-1 SAR data in Norway. *Remote Sensing*,
 252 11(23), 2863. <https://doi.org/10.3390/rs11232863>

253 European Avalanche Warning Services. (2022). *Avalanche size classification standards (classes D1*
 254 *through D5)*. <https://www.avalanches.org/standards/avalanche-size/>.

255 Gatti, M. (2026). *AvalCD* [Dataset]. Zenodo. <https://doi.org/10.5281/zenodo.15863589>

256 Gatti, M., Mariani, A., Gallo, I., & Monti, F. (2026). *Large-scale avalanche mapping from SAR*
 257 *images with deep learning-based change detection*. arXiv preprint arXiv:2603.22658.
 258 <https://doi.org/10.48550/arXiv.2603.22658>

259 Nguyen, H., Dang, K. B., Giang, T. L., Dang, V. B., Dang, N. V., & Dao, M. D. (2025). Exploring
 260 various deep learning models for high-precision landslide tracing in very-high resolution remote
 261 sensing imagery. *Journal of the Indian Society of Remote Sensing*, 53, 1–23.
 262 <https://doi.org/10.1007/s12524-025-02267-z>

263 OpenStreetMap contributors. (2026). *OpenStreetMap*. <https://www.openstreetmap.org/copyright>.

264 Tompkin, C., & Leinss, S. (2021). Backscatter characteristics of snow avalanches for mapping with
265 local resolution weighting. *IEEE Journal of Selected Topics in Applied Earth Observations and*
266 *Remote Sensing*, 14, 4452–4464. <https://doi.org/10.1109/JSTARS.2021.3074418>

267 Wang, S., Wang, D., Shi, J., Zhang, Z., Wang, H., Li, X., & Guo, Y. (2025). Unsupervised SAR
268 image change detection method based on fused difference images. *Journal of the Indian Society of*
269 *Remote Sensing*, 53(12), 4373–4382. <https://doi.org/10.1007/s12524-025-02259-z>

270 Zhang, Z., Liu, Q., & Wang, Y. (2018). Road extraction by deep residual U-Net. *IEEE Geoscience*
271 *and Remote Sensing Letters*, 15(5), 749–753. <https://doi.org/10.1109/LGRS.2018.2802944>